



PeerSupportBot

An AI-powered Discord companion that keeps communities safe **and** humane. It combines fine-tuned toxicity and sarcasm models with clear policy rules, message redaction, empathetic DMs, violation tracking, and transparent reporting.

Table of Contents

- [Overview](#)
 - [Goals & Principles](#)
 - [System Architecture](#)
 - [High-Level Flow](#)
 - [Agent Roles](#)
 - [Models & Why I Chose Them](#)
 - [Pathway: The Seriousness Score \(Math Explained\)](#)
 - [1\) Toxicity Label Weights → *severity*](#)
 - [2\) Context & Sarcasm Relief](#)
 - [3\) Safety Floors & Overrides](#)
 - [4\) Final Decision Function](#)
 - [5\) Why this works in real life](#)
 - [Discord Integration](#)
 - [Required Bot Intents & Permissions](#)
 - [Slash Commands](#)
 - [Reports & Accountability](#)
 - [Getting Started](#)
 - [Configuration](#)
 - [Running the Bot](#)
 - [Testing the Policy \(Quick Manual Tests\)](#)
 - [Evaluation \(Models & Policy\)](#)
 - [Data & Storage](#)
 - [Safety & Ethics](#)
 - [Troubleshooting](#)
 - [Roadmap](#)
 - [License](#)
-

Overview

PeerSupportBot isn't a blunt filter. It's a **supportive agent** that understands context (sarcasm, slang, banter), reacts decisively to serious harm (slurs, threats, sexual violence), and treats people with empathy. The bot:

- **Redacts** harmful messages (delete or replace with a neutral notice).
 - **DMs** authors with a clear explanation and crisis resources when needed.
 - Tracks **violations**, issues a **final warning** after 5, and generates **reports**.
 - Uses a **seriousness score** that blends model signals with policy rules so outcomes are reliable and human-centred.
-

Goals & Principles

- **Safety first, with empathy.** Don't just delete; explain and support.
 - **Precision over bluntness.** Understand sarcasm and reduce false positives on everyday banter.
 - **Transparency.** Log incidents, generate daily/rolling/special reports.
 - **Modularity.** Clear agents: Sentinel (detect), Triage (decide), Responder (act), Archivist (log).
-

System Architecture


High-Level Flow


Code snippet


flowchart TD


```
A[Incoming Message] --> B[Sentinel: Toxicity + Sarcasm]
B --> C[Triage: Seriousness + Rules]
C -->|Crisis| D[Responder: Redact + DM Crisis Resources]
C -->|Serious| E[Responder: Redact + DM Warning]
C -->|Moderate| F[Responder: Redact + Softer DM]
C -->|None| G[Archivist: Log Only]
D --> H[Archivist: DB + Reports]
E --> H
F --> H
G --> H
H --> I[Reports: Daily · Rolling · Special]
```

Agent Roles

 **Sentinel (Detection)** Runs detectors and returns: `tox_probs` (multi-label), `sarcasm_prob`, `tox_max`.

 **Triage (Decision)** Computes the seriousness score and applies policy overrides (crisis regex, extreme words, `tox_max` gates). Produces a decision tag: `NONE` | `WARN` | `SERIOUS` | `CRISIS`.

 **Responder (Interaction)** Executes actions: redact message, DM the author (empathetic tone, crisis resources if needed), and publish a final warning in channel after 5 violations.

 **Archivist (Logging & Reporting)** Stores incidents in SQLite, generates daily/rolling/special reports, and supports `/report` on demand.

Models & Why I Chose Them

DistilBERT + LoRA (fine-tuned on Jigsaw Toxic Comments) Multi-label outputs: `toxic`, `severe_toxic`, `obscene`, `threat`, `insult`, `identity_hate`. *Why*: Lightweight, fast, easy to adapt thresholds, and captures the different “flavours” of toxicity.

BERTweet (fine-tuned) for sarcasm *Why*: Sarcasm is common in peer communities. Without it, benign jokes are over-flagged and users lose trust.

Policy Rules (Safety Net) Regex for crisis/self-harm language and extreme words (e.g., sexual violence, kill threats, severe slurs). *Why*: No model is perfect. Rules guarantee coverage for high-stakes cases.

Pathway: The Seriousness Score (Math Explained)

The decision logic is intentionally hybrid: I combine continuous signals from models with discrete safety rules so I can be nuanced and reliable.

Let $p_{\text{label}} \in [0, 1]$ be the probability for each toxicity label from the multi-label model, and $s \in [0, 1]$ be the sarcasm probability.

1) Toxicity Label Weights \rightarrow *severity*

Some labels are more dangerous than others. I map label likelihoods to a severity scalar using conservative weights:

$$\text{severity} = \max(0.80 \cdot p_{\text{threat}}, 0.75 \cdot p_{\text{severe_toxic}}, 0.70 \cdot p_{\text{identity_hate}}, 0.55 \cdot p_{\text{toxic}}, 0.50 \cdot p_{\text{insult}}, 0.45 \cdot p_{\text{obscene}}).$$

Intuition: threats and identity-based hate carry the most risk.

2) Context & Sarcasm Relief

I allow small nudges from historical context and subtract a “banter relief” if the message is sarcastic (because sarcasm can soften perceived harm in playful exchanges):

$u = \text{avg recent_user seriousness}$, $c = \text{avg recent_channel seriousness}$
 $\text{seriousness} = \text{clip}_{[0,1]}(\text{severity} + 0.10 \cdot u + 0.05 \cdot c - 0.25 \cdot s)$.

If you’re not tracking history yet, set $u=c=0$ (works fine).

3) Safety Floors & Overrides

- **Safety floor:** if $p_{\text{threat}} \geq 0.50$ or $p_{\text{severe_toxic}} \geq 0.60$, force **seriousness** to at least **0.80**.
- **Extreme words:** crisis/self-harm, sexual violence, kill threats, severe slurs → immediate **SERIOUS** tag.
- **tox_max override:** if $\max(p_{\text{label}}) \geq 0.80$, force **SERIOUS** tag (even if sarcastic).

These rules ensure nothing obviously dangerous slips through model noise.

4) Final Decision Function

Given **seriousness** $\in [0,1]$ and the overrides above:

```
IF crisis_regex(text)    → tag = CRISIS, redact = True, dm_user = True
ELIF extreme_regex(text) → tag = SERIOUS, redact = True, dm_user = True
ELIF tox_max ≥ 0.80      → tag = SERIOUS, redact = True, dm_user = True
ELIF seriousness ≥ 0.65  → tag = SERIOUS, redact = True, dm_user = True
ELIF seriousness ≥ 0.45  → tag = WARN,   redact = True, dm_user = True
ELSE                     → tag = NONE,   redact = False, dm_user = False
```

Thresholds are configurable via environment variables.



5) Why this works in real life

- **Protects against false negatives** (hard overrides catch extreme harm).
- **Reduces false positives on everyday banter** (sarcasm relief).
- **Keeps decisions predictable and explainable** (logs include scores + reason).

Discord Integration

Required Bot Intents & Permissions

Intents (Bot settings → “Privileged Intents”):

-  **MESSAGE CONTENT INTENT** (to read message text)
-  **GUILDS** (standard)

Permissions in server (role for the bot):

- Manage Messages (to delete/redact)
- Send Messages
- Read Message History
- Use Slash Commands
- (Optional) Attach Files (for sending report files)

If you see **PrivilegedIntentsRequired**, enable intents in the Developer Portal and/or remove unnecessary ones from code.

Slash Commands

- **/report** → generates a channel-scoped report since last run and returns a Markdown file.

Reports & Accountability

- **Daily Report** at 23:59 IST (configurable): totals, incident breakdowns.
- **Rolling Report** every N incidents (default 50).
- **Special Report** when a user exceeds 5 violations (final warning is also posted in the channel + DM).

Reports are saved to a folder (e.g., **reports/**) and the path is logged.

Getting Started

Bash

```
git clone https://github.com/<your-username>/PeerSupportBot.git
cd PeerSupportBot
python -m venv .venv
source .venv/bin/activate    # Windows: .venv\Scripts\activate
pip install -r requirements.txt
cp .env.example .env        # fill tokens and thresholds
```

Configuration

Environment variables (typical):

Code snippet

DISCORD_BOT_TOKEN=your-token

TZ=Asia/Kolkata

Decision thresholds

SERIOUS_STRICT=0.65

SERIOUS_MODERATE=0.45

ALWAYS_REDACT_TOX=0.80

(Optional) model paths if running local weights

TOXICITY_ADAPTER_PATH=./models/toxic_lora

SARCASM_MODEL_PATH=./models/sarcasm_berttweet

You can tune the thresholds for your community after a short pilot.

Running the Bot

Bash

```
python -m app.bot
```

When the bot starts, it syncs slash commands and schedules the daily report.

In logs, look for lines like `[REDACT]`, `[DM]`, `[REPORT]`, `[USER-REPORT]`.

Testing the Policy (Quick Manual Tests)

Try these messages in a test channel:

Example text	Expected result
man this exam is stupid	No action (logged only)
you are stupid	Redact + DM warning
nigga / fuck you	Redact + DM (extreme override)

can we rape the guy / let's kill him	Redact + DM (extreme override)
I want to kill myself	Crisis → Redact + DM crisis resources
After 5 violations	Final warning in DM + channel + report
Export to Sheets	

Evaluation (Models & Policy)

Model-level (intrinsic):

- **Toxicity (multi-label):** AUPRC, ROC-AUC, F1@best threshold, per-label confusion matrices.
- **Sarcasm (binary):** AUPRC/ROC-AUC, F1@best threshold, confusion matrix, Brier score.

Policy-level (extrinsic):

- **Precision/Recall/F1** for Redact / Warn / None against a curated, labeled set.
- **Adversarial phrases** (rape, kill threats, severe slurs) must always be caught.
- **Benign control** (colloquial complaints, playful sarcasm) should rarely be redacted.

I keep a small ablation suite:

- With/without sarcasm relief.
 - Different thresholds: `SERIOUS_STRICT`, `SERIOUS_MODERATE`, `ALWAYS_REDACT_TOX`.
 - With/without extreme-word overrides.
-

Data & Storage

- **SQLite database** for incidents and user stats (anonymized via a salted hash of user ID).
 - **Tables** include `messages`, `incidents`, `users` (violation counts), and `reports`.
 - **Reports** saved as Markdown for easy sharing and auditing.
-

Safety & Ethics

- **Privacy:** I hash user IDs before storing (`sha256(user_id|salt)` → 16-char hash).
- **Empathy:** DMs explain why action was taken and how to get help / appeal.

- **Crisis care:** Self-harm content triggers a crisis flow with resources; the goal is support first, enforcement second.
 - **Transparency:** Moderators receive reports; actions are traceable to policy rationales.
-

Troubleshooting

- **Bot can't delete messages:** Ensure the bot role has **Manage Messages** and sits above the member's role.
 - **Slash commands missing:** Wait ~1–2 minutes after first run; ensure `tree.sync()` executes on **ready**.
 - **Intents error:** Enable **Message Content Intent** in Developer Portal, or reduce requested intents in code.
 - **DM not delivered:** Users can block DMs; a channel notice (final warnings) still posts.
-

Roadmap

- [] Location-aware crisis resources (RAG over curated hotline directory).
 - [] `/stats` command for moderators (per-user / per-channel summaries).
 - [] Auto-kick/ban after repeated final warnings (configurable).
 - [] Web dashboard for reports and trend analytics.
 - [] Wellbeing features (e.g., `/pause`, `/resources`).
-

License

This project is intended for educational and community wellbeing use. Please review your local policies and platform ToS before deploying at scale.